

BUILDING THE CENTER OF EXCELLENCE FOR POWER PLATFORM

DRIVING REAL INNOVATION IN THE AGENTIC ERA

**WHY BALANCING GOVERNANCE AND ENABLEMENT IS THE SECRET TO
SUSTAINABLE INNOVATION**

WHITE PAPER - SEPTEMBER 2026

NICOLÁS FERNÁNDEZ

Driving Real Innovation in the Agentic Era

Why Balancing Governance and Enablement Is the Secret to Sustainable Innovation

© September 2026 Nicolás Andrés Fernández. All rights reserved.

First edition, September 2026.

ABOUT THE AUTHOR

Nicolás Andrés Fernández specializes in Microsoft Power Platform, Copilot Studio, and Dynamics 365. His contributions to the Microsoft ecosystem have earned him seven consecutive Microsoft Most Valuable Professional (MVP) awards. With more than 17 years of experience, he has designed, developed, and implemented low-code and CRM solutions for organizations across multiple industries and countries. He is the author of the Building the Center of Excellence for Power Platform book series and has extensive experience creating frameworks for adoption, governance, innovation, and sustainable platform management. Nicolás regularly shares his expertise as a speaker at Microsoft and community events, while also supporting initiatives that help organizations and professionals turn technology into sustainable business value.

COLLABORATORS

Technical reviewer: Franco Martino Gagliardi & Wilmer Alcivar

This white paper is independently published by Nicolás Andrés Fernández. It is not official guidance from Microsoft or any other vendor and does not replace official product documentation, licensing terms, security guidance, professional advice, or the organization's own legal, risk, privacy, compliance, and Responsible AI decisions. Microsoft, Power Platform, Microsoft Copilot Studio, Dataverse, Microsoft 365 Copilot, and other product names referenced in this document are trademarks of their respective owners. References to those products are used only for identification and educational discussion.

Sustainable innovation does not emerge from maximizing control or maximizing freedom. It emerges from deliberately balancing governance and enablement, and from using evidence to decide what deserves to scale.

Contents

Prologue	4
Part I: The Innovation Paradox	6
1. Innovation Without Outcomes	7
2. Why the Agentic Era Changes the Equation	9
Part II: Escaping the False Choice.....	12
3. Governance Is Not the Opposite of Innovation	13
4. Licenses Do Not Create Enablement	16
5. The Balance That Makes Innovation Sustainable.....	20
Part III: Operationalizing the Balance	23
6. The Innovation Hub as the Organizational Integrator	24
7. From Idea to Sustainable Value	28
8. Mechanisms That Create the Balance	33
9. Measuring Whether Innovation Is Real.....	36
Innovation With Direction and Momentum.....	40
Continue the Series.....	41
Sources and Further Reading	42

Prologue

Organizations have never had greater access to tools that can turn an idea into something visible. A maker can assemble a workflow in hours. A small team can connect knowledge, actions, and conversational interfaces without waiting for a traditional development cycle. Generative AI can accelerate research, produce options, draft content and reduce the distance between a business problem and a working prototype. In the agentic era, that prototype may not simply answer a question; it may interpret intent, select tools, retrieve information and initiate actions across business systems.

This abundance creates a seductive impression: that the organization is innovating because more people can build more things, more quickly. Ideas submitted, hackathons held, prototypes built, and agents created are easy to count, but they describe activity rather than value. A growing catalog of pilots may coexist with limited adoption, duplicated effort, unclear ownership, unresolved security concerns, and no credible evidence that the work has improved an outcome that matters. The organization becomes active without becoming more capable.

The problem is not experimentation. Experimentation is indispensable when the technology, the use case and the operating implications are still emerging. The problem begins when experimentation becomes the permanent destination. Ideas enter easily, prototypes attract attention and demonstrations create momentary excitement, but few initiatives acquire the sponsorship, evidence, engineering discipline, adoption support or operational ownership required to survive beyond the pilot. Innovation remains visible but unfinished.

THE CENTRAL QUESTION

How can an organization preserve the energy of experimentation while ensuring that innovation remains safe, purposeful, and capable of producing measurable value?

The central thesis of this paper is that sustainable innovation depends on balancing governance and enablement rather than maximizing either one. Governance provides direction through strategic priorities, decision rights, risk boundaries, accountability, and the conditions for scale. Enablement creates momentum through skills, confidence, access to expertise, reusable patterns, communities, and practical support. Evidence adjusts the balance. It determines whether an initiative should continue, adapt, scale, or stop.

Agents make this thesis urgent. They can be created more quickly than conventional solutions, yet their value often depends on access to business knowledge, identity, tools, connectors, and data. Their behavior can vary with context, and their usefulness may change as models, instructions, sources, and business processes evolve. These characteristics increase both the rate of experimentation and the potential consequences of weak design or oversight. An organization that accelerates without proportionate governance accumulates exposure; an organization that controls without building its capacity to learn makes the approved route irrelevant.

Within the Center of Excellence, the Innovation Hub is the organizational integrator that makes this balance operational. It connects strategic intent, governance, development, adoption, community, delivery, and evidence. The Center of Excellence provides the wider framework and mandate. The Innovation Hub provides the execution layer through which worthwhile ideas move from uncertain hypotheses toward trusted, sustainable value.

That movement is only credible when it can be measured. Evidence must show not merely that work is progressing, but whether it is producing value, controlling exposure, and strengthening the organization's capacity to repeat success. Parts I and II distinguish innovation outcomes from activity and establish governance and enablement as complementary forces. Part III translates the balance into an operating model, a pathway, and a practical rhythm. Chapter 9 closes the loop by defining the evidence system leaders need to decide what should continue, adapt, scale, or stop.

The Sustainable Innovation Balance




Figure 1. The Sustainable Innovation Balance: governance as direction, enablement as momentum, evidence as the feedback that adjusts both.

Four implications for leaders



Figure 2. Four implications for leaders.

Part I: The Innovation Paradox

Innovation activity becomes valuable only when learning is converted into a decision.

1. Innovation Without Outcomes

Innovation begins with uncertainty. The organization does not yet know whether an idea will work, whether users will adopt it, whether the economics are attractive or whether the operational consequences are manageable. Experimentation is the disciplined method for reducing that uncertainty. It allows the organization to test a hypothesis before committing to scale.

The distinction matters because an experiment is a means, not an outcome. A proof of concept can demonstrate technical possibility while leaving the business problem untouched. A prototype can create enthusiasm while depending on manual work that will never scale. A hackathon can surface talent and ideas while producing assets with no owner, no support model and no route to production. These activities can be valuable, but they become innovation only when learning is converted into a decision and, where justified, into sustained organizational value.

The experimentation trap

Organizations become trapped when the front door for ideas is easier to design than the path that follows. New initiatives receive attention because novelty is visible. Existing pilots receive less attention because the remaining work, including data quality, architecture, testing, adoption, ownership, and operations, is less glamorous and often distributed across teams. The portfolio grows at the top while the middle becomes congested.

- Prototype theater: demonstrations are treated as evidence of readiness rather than evidence of possibility.
- Solution-first innovation: a new capability is matched to any available problem instead of beginning with a meaningful outcome.
- Duplication by discovery failure: teams rebuild similar agents or components because there is no visible portfolio or reusable catalog.
- Orphaned experimentation: the sponsor funds the pilot but no accountable owner accepts the resulting service.
- Metric substitution: counts of makers, ideas, or agents replace measures of value, reliability, adoption, and learning.

Over time, these patterns create innovation debt. Innovation debt is the accumulated cost of unfinished decisions: prototypes that still consume environments and attention; data access that was granted temporarily but never reviewed; unclear ownership; duplicated assets; and stakeholder expectations that remain unresolved. Like technical debt, it is not inherently wrong. Some uncertainty must be carried while learning occurs. It becomes damaging when the organization cannot see it, value it or retire it.

From activity to value

A useful definition of innovation must therefore include conversion. In this paper, real innovation means the successful conversion of a relevant idea into a measurable improvement that the organization can trust, operate, and sustain. The improvement may increase revenue, reduce cost, improve service, strengthen resilience, reduce risk, increase employee capacity or create a strategic option. What matters is that the intended value is explicit and the evidence is credible enough for the next decision.

WORKING DEFINITION

Innovation becomes real when a relevant idea progresses from possibility to measurable organizational value, and the organization can sustain the result responsibly.

The portfolio must learn

The unit of learning is not only the individual experiment. A mature organization learns across the portfolio. It can see where several teams are testing the same assumption, where a shared dependency is delaying multiple initiatives, and where one experiment has produced an asset or a warning that should change the route for others. Portfolio visibility therefore matters before scale. It prevents the Center of Excellence from treating every request as an isolated conversation and allows the Innovation Hub to compare opportunities against the same strategic themes, risk boundaries, and evidence expectations.

This portfolio view also changes how stopping is interpreted. An experiment may stop because the problem is less important than expected, the proposed approach cannot meet the requirement, the economics are unattractive, the risk cannot be reduced proportionately, or another initiative has already solved the need. Each reason produces different learning. Recording the decision, the evidence, and what can be reused converts a local stop into organizational capability. Without that discipline, the same weak hypothesis may return through another team and consume the same attention again.

Learning must also affect the system around the experiments. Repeated data-quality failures may indicate that a shared data product is needed. Repeated security questions may justify a new approved pattern. Repeated difficulty identifying an owner may expose a sponsorship problem rather than a technical problem. The Innovation Hub should therefore ask two questions after every meaningful decision: what did this initiative learn, and what should the Center of Excellence change because of that learning?

This definition does not demand that every experiment succeed. A well-designed experiment that disproves a hypothesis can protect the organization from a larger investment and can produce reusable learning. The failure is not stopping. The failure is continuing without evidence, or allowing an inconclusive pilot to remain permanently between experimentation and production.

2. Why the Agentic Era Changes the Equation

The balance between governance and enablement has always mattered. Low-code platforms already demonstrated that lowering the barrier to creation expands both opportunity and the need for a coherent operating model. Agents increase the urgency because they combine accelerated creation with access, interpretation, and action.

Microsoft’s 2026 Work Trend Index frames the organizational opportunity around agents and human agency: as agents take on more execution, organizations must redesign how people exercise judgment, direct work and create value. Microsoft’s current guidance for an agentic Center of Excellence makes a related operational point: the Center of Excellence should turn strategy and intent into repeatable, trusted execution, making the safe path the easy path. These positions reflect a wider reality. The limiting factor is no longer simply whether an organization can build an agent. It is whether the organization can repeatedly identify the right work, design the right degree of agency and run the resulting capability responsibly.

Four multipliers

Multiplier	What changes	Why imbalance becomes costly
Speed	Natural-language and low-code experiences shorten the path from idea to prototype.	Poor choices can proliferate before portfolio visibility and standards catch up.
Agency	Agents can select tools, retrieve knowledge and initiate actions within defined permissions.	Weak ownership or controls may affect business processes, data, and customers rather than only a demonstration.
Variability	Responses and decisions depend on instructions, context, knowledge, models, and tool results.	Validation must address quality and behavior over time, not only whether a deterministic workflow executes.
Connectivity	Value often depends on connectors, application programming interfaces, enterprise knowledge and cross-platform actions.	The useful solution crosses organizational boundaries, multiplying dependencies and decision rights.

Table 1. The agentic urgency

These multipliers change both sides of the balance. Governance must become more adaptive, because static documentation and one-time approval cannot capture a system whose behavior and dependencies evolve. Enablement must become more sophisticated, because prompting alone is not enough. Makers and

teams need to understand problem framing, knowledge quality, tool descriptions, identity, security, testing, human oversight, failure modes, operational ownership and value measurement.

The cost of moving too fast

Without adequate governance, organizations can create agents with unclear purpose, excessive access, weak source quality, duplicated capability or no accountable owner. A technically impressive agent may be unsuitable for the decision it influences. An agent may also remain in use after its knowledge, instructions, or process context have changed. The risk therefore exists at creation, at deployment and throughout operation.

Exposure is usually described in organizational terms: processes, data, compliance, and trust. It is worth naming the people behind those terms, because they are the reason oversight exists. A wrong answer reaches a customer who cannot see how it was produced. A redesigned process reaches an employee whose judgment used to be the control. A confident but ungrounded recommendation reaches a colleague who has no practical way to check it. The severity of an error is not only what it costs the organization; it is who absorbs it, and whether that person can recognize and correct it. That question should shape the evidence standard as much as any measure of business impact.

The cost of moving too slowly

Excessive friction creates a different risk. When teams cannot discover the approved route, obtain a suitable environment, access patterns or get timely advice, they stop engaging with the Center of Excellence. Some ideas disappear; others migrate to uncontrolled tools or informal experiments. The organization then loses both the innovation and the visibility required to govern it. In the agentic era, the goal is a trusted route from idea to responsible operation.

Urgency does not justify indiscriminate acceleration

The urgency of the agentic era can itself distort decision-making. Leaders may interpret rapid market change as a reason to bypass problem definition, compress validation, or accept unclear ownership. That response creates speed at the beginning and delay later, when the organization discovers that the agent depends on weak knowledge, cannot be supported, or has no credible measure of value. Responsible urgency means reducing avoidable delay while preserving the learning required for the next commitment.

The distinction between reversible and difficult-to-reverse decisions is useful here. A small, time-limited experiment with synthetic or low-sensitivity data may justify rapid approval because exposure is constrained and the organization can stop easily. A decision to integrate an agent into a critical process, expose it to a broad audience, or allow it to perform consequential actions requires stronger evidence because reversal may affect customers, operations, trust, and compliance. The objective is not to move every initiative at the same speed. It is to make the cost of delay and the cost of error visible at the same time.

This is why agentic innovation needs both operating tempo and decision discipline. Tempo keeps the approved route relevant. Decision discipline prevents urgency from becoming an excuse for unexamined

commitments. The balance between governance and enablement allows the organization to vary both according to the decision in front of it.

Part II: Escaping the False Choice

Governance creates direction. Enablement creates the confidence and capability to move.

3. Governance Is Not the Opposite of Innovation

Governance is often experienced through its most visible mechanisms: policies, approvals, restrictions, and reviews. When these mechanisms are disconnected from purpose, slow to operate or insensitive to context, governance appears to be the institutional opponent of innovation. Teams learn that the fastest route is to minimize visibility until the work is too advanced to stop.

That is not an argument against governance. It is evidence of governance that has been designed as administration rather than as an innovation system. Effective governance answers practical questions early: What kinds of problems are strategically relevant? Who may experiment, and where? What data or actions require additional scrutiny? What evidence is needed before broader use? Who becomes accountable if the experiment succeeds? What conditions trigger review, restriction, or retirement?

Governance as a design discipline




Figure 3. Governance as a design discipline.

Proportionality is especially important. The earliest experiment should be constrained enough to learn safely, but it should not carry the full burden of a production service. Conversely, an agent that influences regulated decisions, accesses sensitive information or performs high-impact actions should not inherit the lightweight treatment of a personal productivity experiment. A single universal process is easy to document and difficult to defend.

Proportionality only becomes usable when the organization can say what the tiers are. The following illustration shows the shape rather than the answer: the dimensions that set a tier, and what changes as the tier rises.

Tier	What places an initiative here	What changes
Contained	Synthetic or low-sensitivity data, no consequential action, a named small audience, easily reversible.	A shared experimentation environment, self-service approval, evidence limited to whether the problem is real.

Tier	What places an initiative here	What changes
Supported	Business data, actions that write to systems of record, a defined department, reversible with effort.	A dedicated environment, an accountable owner before validation, test evidence across representative scenarios, a named reviewer.
Assured	Sensitive or regulated data, consequential or customer-facing actions, broad reach, difficult to reverse.	Production controls from the outset, security, and responsible AI assessment, explicit residual-risk acceptance, monitoring, and recertification.

Table 2. An illustrative risk classification

This table is an example, not a standard. The dimensions are durable; the number of tiers, their names and their thresholds should be set by each organization and revised as evidence accumulates. A classification copied unchanged from a document is a classification nobody has thought about.

A tiering model of this kind is one way of doing what published risk frameworks describe in the abstract. The generative AI profile of the NIST AI Risk Management Framework organizes the work into governing, mapping, measuring, and managing. Governing is the decision rights and the appetite this paper places above the Innovation Hub. Mapping is the classification itself, the point at which context and consequence are established. Measuring is the evidence discussed in chapter 5 and the scorecard in chapter 9. Managing is the operating rhythm, the transitions and the retirement mechanisms. The OECD principles operate at a different level, describing the qualities a system should have rather than the process that produces them, and they are most useful here as a test applied at the mapping stage. Naming a framework is easy. The value comes from being able to point at which of its functions an operating model actually performs, and which it does not yet.

Guardrails should create a usable path

Copilot Studio and Power Platform provide governance capabilities such as environments, roles, data policies, connector controls, publishing controls, application lifecycle management and administrative monitoring. These controls matter, but technology alone does not create governance. The organization must translate policy into a coherent route that people can understand and follow. A technically secure environment with an opaque process still encourages circumvention.

The most useful governance mechanisms are those that reduce repeated negotiation. A risk classification determines the pathway. An approved pattern resolves a common architecture question. A reusable connector encapsulates a safe integration. A decision record explains why a boundary exists. A service-level expectation tells teams when they will receive an answer. In each case, governance increases speed by making a recurring decision once and making the result visible.

Identity deserves particular attention, because it is where the agentic change is sharpest. A conventional application acts as the person using it, and the permissions question is answered by that person. An agent may act on behalf of a user, under its own identity, or through a service principal whose permissions were granted for a different purpose entirely. Whatever it can reach, it can reach on behalf of whoever invokes it, and it decides for itself which of those reaches to use. Least privilege is therefore harder to reason about, not because the controls are different but because the actor is not deterministic: the set of resources an agent will touch cannot be enumerated in advance from its configuration alone. Governance should establish which identity an agent runs as, what it inherits from that identity, which connectors and knowledge sources are in scope for that tier, and how an inherited permission is reviewed when the process around the agent changes.

Governance must reduce decision latency

Teams experience governance not only through the quality of its rules but through the time and uncertainty involved in obtaining a decision. A correct control that arrives after the experiment has lost sponsorship, missed its business window, or moved to an informal route has not supported responsible innovation. Decision latency should therefore be treated as a governance design measure, not merely as an administrative service level.

Reducing latency does not mean approving more work. It means making more decisions predictable. The organization should distinguish questions that can be answered by a published pathway, questions that require a specialist, and questions that require accountable risk acceptance. It should also make the evidence request proportionate and explain what happens next. Teams can plan around a demanding standard when the route is visible; they cannot plan around silent queues, changing expectations, or reviews that restart the argument from the beginning.

The Innovation Hub can expose recurring decision patterns and help the Center of Excellence convert them into reusable policy, architecture, and enablement assets. Over time, fewer initiatives should need bespoke interpretation. Specialists can then focus on genuinely novel or consequential questions, while low-risk work progresses through supported self-service. This is one of the clearest signs that governance is maturing from review activity into an organizational capability.

DESIGN PRINCIPLE

Good governance defines the conditions under which innovation can progress responsibly. The practical test is whether it helps a capable team move a worthwhile idea forward.

4. Licenses Do Not Create Enablement

A license creates access to a product. It does not create judgment, confidence, trust, community, operating discipline or the ability to turn a useful idea into a sustainable outcome. Treating license allocation as enablement confuses technical availability with organizational readiness.

Training transfers knowledge at a point in time. Enablement creates the conditions in which people can repeatedly apply knowledge, obtain help, collaborate across disciplines and produce outcomes within the intended operating model. The distinction is critical in an environment where platforms, models, and guidance change faster than a static curriculum can be refreshed.

An organization can distribute thousands of licenses, train hundreds of makers and still fail to innovate. People may understand the interface but lack access to a safe environment, meaningful data, reusable patterns, expert advice, a sponsor or a route to production. Access can increase demand while the rest of the system remains unable to absorb it. The result is frustration, scattered experimentation and false confidence rather than capability.

The enablement system

Capability	What it provides	Innovation effect
Orientation	Clear entry points, use-case guidance and a visible route from idea to operation.	Reduces wasted effort and solution-first experimentation.
Skills	Role-based learning across business, architecture, security, responsible AI and operations.	Improves the quality of decisions, not only build speed.
Patterns	Approved architectures, templates, components, test approaches and examples.	Makes good practice reusable and lowers the cost of compliance.
Expert access	Coaching, office hours, reviews, and specialist participation at the right stage.	Resolves uncertainty before it becomes rework.
Community	Peer learning, champions, showcases, feedback, and recognition.	Distributes knowledge and sustains momentum beyond formal programs.
Delivery support	Fusion-team collaboration and a clear escalation route for complex or critical solutions.	Matches the delivery model to the problem rather than forcing every team into self-service.

Table 3. Enablement as organizational capability

Learning to work with AI

Enablement and adoption must now include a new form of professional literacy: learning how to work with AI. People need to understand when to rely on an agent, when to question it, how to verify what matters and when human judgment must remain decisive. Appropriate trust is neither blind acceptance nor automatic rejection. It is confidence calibrated through evidence, context, and consequence.

Teams also need to learn how much risk they are prepared to take. That means identifying what could go wrong, assessing the severity and reversibility of the outcome, choosing safeguards, defining human oversight and deciding what evidence is sufficient for the next level of exposure. This is a practical capability, not a compliance lecture. It allows people to move faster because uncertainty is named and managed rather than ignored.

Capabilities for working responsibly with AI




Figure 4. Capabilities for working responsibly with AI.

The Community Hub makes enablement visible

As I describe in *Designing the Community Hub*, the Community Hub becomes valuable when it makes the Center of Excellence visible, accessible, and actionable in everyday work. It is not another intranet destination. It is the community-facing surface where people discover documentation, approved patterns, frequently asked questions, decision guides, events, champions, reusable assets and routes to help.

Its value increases when it is integrated with the Innovation Hub. Ideas and feedback should enter the Innovation Hub through a visible idea funnel. Training requests should reach the adoption pillar. Documentation and reusable assets should reflect current governance and delivery learning. Community conversations should expose repeated friction and emerging opportunities. The Community Hub makes participation easy; the Innovation Hub converts that participation into prioritized action, supported experimentation and portfolio decisions.

Events, hackathons, and showcases generate energy. They create sustainable value only when participants can find standards, reuse assets, obtain expertise, submit ideas and progress successful work through a coherent operating model. Community is therefore not a communications layer added after the operating model is designed. It is part of how the operating model learns.

Champions are the part of that system most often reduced to a title. A champion network works when it has a lifecycle rather than a list: people are identified through the work they already do, given early access, materials, and a direct line to the Center of Excellence, asked for something specific such as running a clinic

or reviewing a pattern, recognized in a way their manager can see, and released when the role no longer fits their work. Without that cycle the network ages quietly, and the organization keeps naming champions who stopped participating some time ago.

Adoption means changing work, not attending training

Training can explain a feature, a pattern, or a policy. Adoption is demonstrated when people change how work is performed and can sustain that change. For agentic solutions, this includes learning which tasks should be delegated, what context must be supplied, how outputs should be reviewed, when escalation is required, and how responsibility remains with the human and organizational roles around the agent. A completion record from a course says little about whether those behaviors have become part of everyday work.

Enablement must therefore connect learning to real use cases. Role-based scenarios should allow people to practice judgment in the conditions they will actually face. Managers need to learn how to redesign work and set expectations. Makers need to learn how purpose, instructions, knowledge, tools, identity, and testing interact. Reviewers need to understand how evidence changes with risk and intended reliance. End users need to recognize the limits of the experience and know how to challenge, correct, or escalate. These are different learning needs, even when everyone uses the same platform.

Asking users to challenge an agent places a duty on them, and the duty is only fair if the experience supports it. People can only question what they can see. The experience should make plain that they are working with an agent rather than a person, show which sources an answer was drawn from, and indicate where the agent is operating outside what it can support. Correction needs somewhere to go: a challenge that ends in a private conversation improves nothing, whereas one recorded against the agent, its knowledge source or its instructions becomes evidence that changes the system. Escalation needs a destination and a response time, because an escalation route that nobody answers teaches people to stop using it.

The Community Hub is the durable surface for this learning system. It can connect formal learning with decision guides, examples, approved patterns, office hours, champions, frequently asked questions, and the idea funnel. More importantly, it can keep those assets synchronized with governance and delivery experience. When a recurring error appears in validation, the response should not remain inside the project team. It should become a clearer example, an updated checklist, a new test pattern, or a community conversation that improves future work.

Adoption also requires leaders to make room for learning. People cannot develop calibrated trust if they are rewarded only for speed or instructed to use AI without time to verify it. A sustainable approach recognizes an initial learning curve, makes responsible challenge legitimate, and measures whether the new way of working improves the intended outcome. The objective is not usage for its own sake. It is competent reliance where the evidence justifies it.

Enablement makes governance usable

A policy can say that sensitive data must not be exposed. Enablement explains how to classify the use case, which environment to use, which connectors are approved, how to test access boundaries and where to obtain advice. A policy can require an owner. Enablement provides a sponsor canvas, operating-model checklist and examples of ownership responsibilities. In this sense, enablement does not oppose governance; it makes governance usable.

Enablement also improves governance itself. Communities expose recurring friction. Office hours reveal ambiguous policies. Reviews identify missing patterns. Makers discover use cases that central teams did not anticipate. A mature Center of Excellence listens to these signals and changes the system. Governance that never learns from enablement becomes detached from practice; enablement that never feeds back into governance becomes an informal support channel without structural influence.

DESIGN PRINCIPLE

Enablement is the organizational capability to turn intent into competent action within a trusted operating model.

5. The Balance That Makes Innovation Sustainable

The Sustainable Innovation Balance introduced in the Prologue is not a recommendation to place governance and enablement in separate columns and give them equal budget. It is a dynamic operating principle. The appropriate balance changes with the maturity of the organization, the stage of the initiative, the sensitivity of the data, the autonomy of the agent, the intended audience and the evidence accumulated so far.

Governance provides direction

Direction includes strategy, boundaries, and accountability. It helps the organization decide which opportunities deserve attention, what must remain protected and who holds the decision rights. Without direction, enablement increases the number of possible journeys without improving the probability that they reach a worthwhile destination.

Enablement creates momentum

Momentum includes capability, confidence, participation, and access to support. It allows teams to move through the defined route and to make better decisions without escalating every question. Without momentum, governance may describe a safe destination but provide no practical way to reach it.

Evidence adjusts the balance

Evidence is the connective mechanism. Early evidence may be qualitative: a confirmed problem, user interviews, a process baseline or a credible sponsor. Validation evidence may include test results, observed behavior, security findings, user feedback and economic assumptions. At scale, evidence includes adoption, reliability, cost, quality, risk events, realized benefits and operational learning. The evidence required should be proportionate, but a transition without evidence is merely optimism expressed as governance.

What counts as sufficient evidence

Evidence is sufficient when it reduces the uncertainty that matters for the next decision. It does not need to answer every future question. During discovery, the organization may need to know that the problem is real, material, and connected to a strategic priority. During validation, it may need to know that the agent performs acceptably across representative scenarios, that users understand its role, and that material risks have defined treatments. Before production, it needs evidence that ownership, support, monitoring, security, adoption, cost, and value can be sustained at the intended reach.

Four qualities make evidence decision-worthy. First, it must be relevant to the claim being tested. A successful demonstration does not prove adoption, and positive user feedback does not prove security. Second, it must be credible enough for the consequence of error. The organization should distinguish observation, validated estimate, expert judgment, and untested assumption instead of blending them into one business case. Third, it must be current. Evidence can expire when the model, knowledge source,

process, user population, or operating context changes. Fourth, it must be owned. Someone must be accountable for the measure, the interpretation, and the decision it informs. A fifth quality applies wherever an agent influences how people are treated: the evidence should say for whom the system works and for whom it works less well. Aggregate performance can conceal a group for which the experience is materially worse, and an average that looks acceptable is not evidence that the outcome is acceptable to everyone it reaches. Where an initiative affects access, eligibility, prioritization, or service quality, the evidence should be capable of showing that difference rather than averaging it away.

The standard should become stronger as reliance increases. Early evidence can be directional because the organization is deciding whether to learn more. Production evidence must support a decision to expose users and processes to an ongoing service. High-impact actions, sensitive data, regulated decisions, or limited reversibility increase the required depth. Low-risk, reversible experiments should not be forced to satisfy the same burden before they can begin.

Evidence must also be capable of changing the answer. If every initiative progresses regardless of the findings, evidence is decoration. The Innovation Hub should define in advance which results would support continuation, adaptation, scale, or stop. This makes the experiment intellectually honest and protects the decision from novelty, hierarchy, and sunk-cost pressure. It also creates a record that can improve future pathways and standards.

Imbalance	Observable pattern	Likely consequence
Governance dominates	Long queues, universal reviews, centralized decisions and unclear self-service pathways.	Slow learning, shadow experimentation and a Center of Excellence perceived as a gatekeeper.
Enablement dominates	High activity, broad access and strong community energy with weak portfolio discipline.	Duplication, unmanaged exposure, orphaned agents and limited value evidence.
Evidence is weak	Initiatives progress through enthusiasm, seniority, or technical novelty.	Investment persists without learning; stopping becomes politically difficult.
Balance is deliberate	Risk-based pathways, reusable support and explicit evidence at transitions.	Faster responsible learning and a higher probability that successful ideas can scale.

Table 4. Four operating states

Balance is not compromise

The objective is not to make every stakeholder equally satisfied or to split the difference between speed and safety. Some initiatives should be stopped immediately because they conflict with policy, lack a meaningful owner or cannot create sufficient value. Others should move quickly because the risk is low and

the learning is time-sensitive. Balance means applying the right intensity of governance and the right type of enablement to the decision in front of the organization.

CORE MODEL

Governance establishes the safe direction. Enablement creates the capability and momentum to move. Evidence determines whether the organization should continue, adapt, scale, or stop.

Part III: Operationalizing the Balance

*The Innovation Hub turns balance into an operating system
for sustainable value.*

6. The Innovation Hub as the Organizational Integrator

A Center of Excellence is often described by its components: governance processes, administrators, architects, training, community, standards, and tools. In the framework I present in my books, the Center of Excellence provides the wider strategic and operating model, while the Innovation Hub is the execution layer that integrates governance, development, and adoption around real opportunities.

The distinction matters. Calling the entire Center of Excellence the integrator obscures where work is actually coordinated. The Innovation Hub is the point of convergence inside the Center of Excellence. It translates strategy into an idea portfolio, joins ideas with proportionate governance and delivery support, connects adoption and community to execution, and returns evidence to the wider operating model.

The Innovation Hub does not act alone, and it is not one of the centers. The Center of Governance owns boundaries, assurance, and the conditions under which work may progress. The Center of Enablement owns skills, community, learning, and adoption, and the Community Hub is its community-facing surface: the center maintains the documentation, the patterns, and the decision guides that the Hub makes discoverable. The Center of Innovation defines the development practices and delivery models, and it leads the Innovation Hub, holding the largest role in it. The Hub itself is transversal: it runs across the three centers as a coordination layer, ensuring balance and the fluidity of the processes that move between them, and it is how the outcomes the Center of Excellence pursues are made real. Sustain runs underneath all three, keeping alive what they produce.

Microsoft's 2026 guidance for an agentic Center of Excellence similarly positions the wider Center of Excellence as an organizational vehicle that turns strategy and intent into trusted execution. The model in this paper makes the execution mechanism explicit: the Innovation Hub coordinates the movement of ideas, decisions, expertise, and learning. It is materially different from a committee that reviews work after it has been designed or a central team that attempts to build every solution itself.

Something must sit above the Innovation Hub, because the Hub translates a strategy it does not write. Innovation themes, the investment envelope and the risk appetite that governance boundaries express are strategic decisions, and they need an author: a sponsor and a steering forum operating on an annual or half-yearly cadence, distinct from the weekly and monthly rhythm of the Hub. That forum decides which themes are open, how much the organization is prepared to invest in learning, and how much exposure it will accept in exchange. The Innovation Hub then decides, initiative by initiative, what to do within those boundaries and returns the evidence that should change them at the next cycle. Without that layer, the direction the paper attributes to governance has no source, and a team can be told that its idea is off-theme with nobody able to say who set the theme or how it could be revisited.

Five connected responsibilities

Five connected responsibilities




Figure 5. The Innovation Hub as the integrating execution layer, and its five connected responsibilities.

The Innovation Hub should not own every innovation. Business leaders remain accountable for business outcomes. Platform teams remain accountable for platform health and control implementation. Security, compliance, legal, data, and responsible AI functions retain their mandates. Delivery teams remain accountable for design and engineering quality. The Innovation Hub creates the common operating rhythm and removes the gaps between these responsibilities.

Decision rights before process

Many innovation processes fail because they describe activities without clarifying authority. A team may complete an assessment, architecture review and pilot, yet no one knows who can authorize production, accept residual risk, fund operation or stop an underperforming service. A useful Innovation Hub therefore defines decision rights around a small set of transitions and makes the accountable role visible.

Decision	Accountability question	Typical Innovation Hub contribution
Enter validation	Is the problem relevant and the hypothesis worth testing?	Problem framing, duplication check, initial risk route and evidence criteria.
Expand exposure	Is there sufficient evidence to involve more users, data, or actions?	Validation review, control requirements and updated ownership.
Enter production	Can the organization operate, support, and govern this capability?	Readiness criteria, architecture assurance and operational handover.

Decision	Accountability question	Typical Innovation Hub contribution
Continue or retire	Is the agent still valuable, trusted, owned, and relevant?	Portfolio evidence, recertification triggers and retirement support.

Table 5. Decisions the operating model must make explicit

Of these, the handover into production is the one most often assumed rather than designed, and it is the transition that decides whether orphaned experimentation happens again. A handover transfers more than a working solution. It transfers the service and its users, the outstanding defects and improvements, the cost of running it, the accepted residual risk, and the monitoring that will reveal when any of those change. The receiving owner is entitled to refuse what is not ready, and is entitled to ask for the evidence on which that judgment rests: what breaks, how often, who has been answering questions so far, what the knowledge sources depend on, and what the agent will cost to run at the intended reach. A transfer that cannot answer those questions has not been arranged; it has been announced.

Community Hub and Innovation Hub: different roles, one system

The Community Hub and the Innovation Hub are closely connected, but they are not interchangeable. The Community Hub is the accessible surface through which the wider organization discovers the Center of Excellence, finds documentation and approved assets, joins the community, requests help, and enters the idea funnel. The Innovation Hub is the integrating execution layer that receives those signals and coordinates shaping, prioritization, risk routing, delivery support, evidence, and portfolio decisions.

This distinction protects both functions. If the Community Hub attempts to own portfolio decisions, it becomes more complex and less accessible. If the Innovation Hub becomes the only place where knowledge and support can be found, participation depends on personal contacts and the operating model becomes invisible. One function optimizes discoverability and participation; the other optimizes conversion and decision-making. Their integration allows a person to move from curiosity to a supported route without needing to understand the internal structure of the Center of Excellence.

The connection must work in both directions. Ideas, support requests, and community feedback flow from the Community Hub into the appropriate Center of Excellence function. Decisions, standards, reusable assets, learning, status, and recognized examples flow back to the Community Hub. The idea funnel is one visible expression of this connection, not the whole relationship. The Community Hub is the front door and shared knowledge surface. The Innovation Hub is the operating engine that turns selected demand into coordinated action and returns learning to the organization.

The precise structure may be centralized, federated, or hybrid. The key test is whether the organization experiences one coherent system. If governance, community, architecture, platform administration and value management operate as unrelated programs, the Center of Excellence exists in name but the Innovation Hub is not functioning as its integrator.

INNOVATION HUB PROPOSITION

The Innovation Hub is not the owner of every idea. It is the integrating execution layer that allows worthwhile ideas to become trusted, sustainable value inside the Center of Excellence.

7. From Idea to Sustainable Value

A sustainable innovation pathway should reduce uncertainty in stages. Each stage has a different purpose and therefore requires a different balance of governance, enablement, and evidence. The pathway is not intended to create a rigid funnel in which every idea follows identical steps. It creates a common language for deciding what the organization needs to learn next. That distinction should be read before the stages themselves. Chapter 3 argued that a single universal process is easy to document and difficult to defend, and nothing here contradicts it: the five stages are a decision vocabulary, not a mandatory sequence, and their purpose is to make the transitions legible rather than to make every initiative travel the same distance.

The shape will be familiar. Staged investment with evidence at each gate is older than any technology this paper discusses, and the discipline of stating a hypothesis, running the smallest useful test and letting the result decide comes from lean and iterative practice. What is adapted here is the emphasis. The stages are not phases of delivery, and they say nothing about how the work is organized inside them; a team may run the whole of Shape and Validate in a single agile cadence. They exist to name the moments at which the organization commits further exposure, and to make the evidence for that commitment explicit.

The idea funnel is managed through the Innovation Hub and made accessible through the Community Hub. The Community Hub gives people a visible place to understand what constitutes an idea, discover existing solutions, review examples and submit feedback. The Innovation Hub then owns the acknowledgement, shaping, prioritization, route, and decision cadence. This separation keeps participation simple without reducing portfolio management to an unstructured suggestion box.

From idea to sustainable value




Figure 6. The five-stage pathway, with evidence-based transitions between stages and a feedback loop from Sustain to the operating model.

Discover: find a meaningful problem

Discovery begins with business friction, unmet need, strategic opportunity or a new capability that may change what is possible. Governance contributes innovation themes, ethical and policy boundaries, portfolio visibility, and a mechanism to detect duplication. Enablement contributes ideation methods, examples, accessible brief templates and contact with people who understand both the business and the platform. The evidence is not a finished business case; it is a credible problem, an affected population and a reason the opportunity matters now.

Shape: turn the idea into a testable hypothesis

Shaping clarifies the intended outcome, sponsor, owner, users, data, actions, constraints, and smallest useful experiment. The team selects the delivery route and identifies the uncertainties that could invalidate the idea. Governance establishes the initial risk tier and non-negotiable boundaries. Enablement provides coaching, patterns, architecture input and access to safe experimentation. The exit question is whether the organization knows what it is trying to prove and can test it proportionately.

Validate: learn under controlled conditions

Validation tests the hypothesis with realistic users, data, and operating conditions while limiting exposure. For an agent, this includes task quality, source grounding, action behavior, identity, failure handling, human oversight, security, and user trust. Governance defines the permitted scope and evidence requirements. Enablement helps the team design test cases, interpret findings and resolve defects. Validation ends with a decision, not merely a demonstration.

Scale: industrialize what works

Scale is a change of state. The concern moves from whether an idea can work to whether the organization can rely on it. Architecture, lifecycle management, support, service ownership, monitoring, adoption, capacity, cost, and change management become central. Governance establishes production readiness and residual-risk decisions. Enablement connects the team to production patterns, operational specialists and adoption support. Evidence must justify the investment and show that the intended outcome remains achievable at broader reach. Cost deserves a specific caution here, because agents invert a familiar assumption. Traditional software rewards scale: the build is paid for once and the unit cost falls as use rises. Consumption-based agents behave in the opposite direction, since every interaction carries a cost and the most successful agent is also the most expensive one to keep running. A benefit case that assumes the usual economics will be wrong in the direction that hurts, and it will be wrong precisely when the initiative succeeds. Cost strategy, licensing models and financial operations are substantial subjects in their own right and belong in the Center of Excellence strategy; they fall outside the focus of this paper. What belongs here is only the consequence for the decision at hand: run cost at the intended reach is part of the evidence for scaling, not a question deferred until after it.

Sustain: keep value, risk, and relevance visible

A production agent is not the end of innovation. It is a live capability whose value and risk can change. Sources become stale, processes change, models evolve, ownership moves and user behavior reveals new needs. Governance provides monitoring, recertification, incident, change, and retirement mechanisms. Enablement provides feedback channels, continuous enhancement and knowledge sharing. The organization must be able to improve, restrict, or retire deliberately.

Sustain is not a stage the Innovation Hub retains. It is a function that receives the capability, and the two work closely: the Hub decides what should be sustained and returns operational learning to standards and

priorities, while sustain holds what is already live and tells the Hub what running it actually costs and where it strains. The handover is only real when the receiving owner, the support route and the funding of operation are agreed before production rather than negotiated after it.

Five mechanisms carry that work, and each needs a trigger, an owner and an outcome rather than a name.

Monitoring is continuous and owned by the service owner. It watches quality, behavior, usage, cost, and the health of connections and knowledge sources, and it exists to make the other four mechanisms fire on evidence rather than on complaint.

Recertification is periodic, and the interval follows the risk tier: frequent for an assured agent, occasional for a contained one. It is the mechanism that detects what otherwise goes unnoticed, since knowledge sources go stale, processes change and models evolve without anyone raising an incident. It asks whether the purpose still holds, whether the sources are current, whether the permissions are still the right ones and whether the owner is still the owner. Its output is what makes stale knowledge measurable at all.

Incident management is triggered by a failure or a harm and owned jointly by support and the service owner. It restores service, records what happened, and asks whether the same weakness exists in comparable agents.

Change management is triggered by any alteration to instructions, knowledge, tools, permissions, or the surrounding process. It matters more for agents than for conventional software, because a change to a source or an instruction can alter behavior without altering anything a release note would capture.

Retirement is triggered by recertification, by a portfolio decision or by the disappearance of the need. It covers withdrawal of access, disposal, or migration of data, communication to users and the release of the cost. An agent that stops being used but is never retired continues to hold permissions and to appear in inventories, which is innovation debt in its most literal form.

Support is part of sustaining, not part of enabling

Office hours, coaching, and clinics help people who are building. They do not help a business user whose agent has just given a confident wrong answer, or a maker whose production solution has stopped working. Those are support, and support is a different discipline.

Two routes are needed and they are not the same. End-user support answers people who are using a capability: it triages between a defect, a gap in training and a request for change, and it needs a response expectation the user can rely on. Maker support answers people who are building and are stuck: it unblocks rather than takes over, and its hardest judgment is when to escalate to a specialist instead of solving the problem on the maker's behalf. Both routes should exist and be published before an agent reaches production. Deciding who answers after the first failure is how a successful pilot becomes an abandoned service.

The transitions matter more than the stages

The five stages provide a common language, but the quality of the operating model is revealed at the transitions. A stage name can create the appearance of discipline even when the initiative moves forward because of enthusiasm, senior sponsorship, or calendar pressure. A transition is meaningful only when the decision, accountable role, required evidence, and resulting conditions are explicit.

Each transition should answer four questions. What uncertainty has been reduced? What new exposure will be introduced if the initiative proceeds? Who is accountable for accepting that exposure and funding the next commitment? What would cause the initiative to pause, change route, or stop? These questions prevent a stage gate from becoming a ceremonial checkpoint. They also allow the organization to use different pathways while preserving a consistent decision logic.

Fast pathways remain possible. A low-risk opportunity may move from discovery to validation in one working session when the problem, owner, boundaries, and test plan are clear. A consequential agent may require several iterations within shaping or validation before broader exposure is responsible. The common model does not force equal duration. It ensures that speed reflects clarity and evidence rather than the ability to bypass scrutiny.

Stage	Governance question	Enablement question	Evidence for the next decision
Discover	Does this fit our purpose and boundaries?	Can people articulate and submit the opportunity?	Problem relevance, affected users and strategic connection.
Shape	What pathway and limits are proportionate?	Can the team design the smallest useful test?	Owner, hypothesis, risk tier and validation plan.
Validate	Are exposure and controls appropriate to learning?	Can the team test quality, behavior, and usability?	Test results, user feedback, risk findings and revised value case.
Scale	Is the capability ready for broader reliance?	Can delivery and adoption be industrialized?	Readiness, operating model, investment case and acceptance.
Sustain	Should the capability continue in its current form?	Can learning drive improvement and reuse?	Value, quality, cost, risk, ownership, and relevance trends.

Table 6. The balance across the lifecycle

The stages and the decisions in chapter 6 describe the same lifecycle from two angles, and they should be read together rather than as two models. The decisions are the transitions between the stages: entering

validation is the move from Shape to Validate, expanding exposure is the move within or beyond Validate as more users, data, or actions come into scope, entering production is the move from Scale into live operation, and continue or retire is the recurring decision inside Sustain. A stage describes what the organization is trying to learn. A decision describes what it is prepared to commit next, and who is accountable for that commitment.

8. Mechanisms That Create the Balance

The balance is created through mechanisms that reduce ambiguity and connect people with decisions. It is not created by publishing a governance policy and a training calendar. The following mechanisms form a minimum viable system; organizations should adapt their form and intensity to context and maturity.

Governance and enablement mechanisms

Mechanisms that create the balance




Figure 7. Governance and enablement mechanisms, paired: every governance mechanism has an enablement counterpart that makes it usable.

The diagram should be read horizontally. Every governance mechanism has an enablement counterpart that makes it usable, and every enablement mechanism needs a governance counterpart that keeps it aligned, proportionate, and sustainable. The Innovation Hub coordinates these pairs; the Community Hub exposes the routes, assets, and learning to the wider organization.

Connecting mechanisms

The most important mechanisms are those that make governance and enablement part of the same experience.

Mechanism	Governance contribution	Enablement contribution
Transparent intake	Captures ownership, value, data, and risk signals.	Gives teams a clear start and rapid orientation.
Experiment charter	Defines scope, boundaries, timebox, and decision rights.	Clarifies the hypothesis, support, and learning plan.
Reusable pathways	Applies consistent controls to recurring risk patterns.	Reduces negotiation and shows teams how to progress.
Evidence-based transitions	Prevents scale without proportionate assurance.	Focuses team effort on what must be learned next.
Portfolio reviews	Aligns investment, risk, and lifecycle decisions.	Connects teams to sponsors, specialists, and reusable learning.
Feedback into standards	Corrects policies that do not match observed reality.	Turns delivery experience into better assets and guidance.

Table 7. Mechanisms that connect direction with momentum

Intake deserves more than a row, because it is where decision latency is created. A transparent intake asks a small number of things and asks them once: what problem this addresses and for whom, who is sponsoring it, what data and actions it would need, whether anything similar already exists, and what would count as success. It should be short enough that a busy person completes it and rich enough that the first routing decision can be made without a meeting. What matters as much as the questions is the promise attached to them. The submitter should know when they will hear back, from whom, and what the possible answers are, including the answer that the idea will not proceed and why. An intake that captures signals without committing to a response is a suggestion box, and people stop using suggestion boxes.

A minimum viable operating rhythm

The operating rhythm should be frequent enough to support learning and light enough to sustain. A practical starting point is a weekly triage for new opportunities and blockers; scheduled specialist clinics; a fortnightly validation and production-readiness forum for initiatives that reach those decisions; a monthly portfolio review covering value, risk, and lifecycle; and a quarterly review of innovation themes, standards, and enablement priorities.

The rhythm is more important than a large initial structure. A small Center of Excellence with an active Innovation Hub, explicit decisions, visible pathways and regular feedback can create more value than a broad committee that meets infrequently and owns no outcomes. As demand grows, automation should remove repeated administrative work and specialists should be engaged according to triggers rather than attending every conversation. That rhythm has a resourcing implication worth stating plainly, since a weekly triage, scheduled clinics, a fortnightly readiness forum, a monthly portfolio review and a quarterly theme review together describe a staffed function rather than a distribution of goodwill. The organization does not

need to resolve how the Innovation Hub is funded before starting, but it should recognize that an unfunded coordination layer degrades into a meeting that people stop attending, and that the cost of the Hub is small next to the cost of the duplication and abandoned pilots it exists to prevent.

Applying the model at different maturity levels

An early Center of Excellence does not need to implement every mechanism at full scale. It needs a coherent minimum. That minimum includes a visible entry point, a small set of strategic themes or eligibility rules, an initial risk classification, named decision owners, access to guided experimentation, and a recurring review that can continue, adapt, scale, or stop work. The first goal is not automation. It is consistency and visibility.

As demand increases, repeated decisions should become pathways, templates, and approved assets. The Community Hub should make them discoverable. The Innovation Hub should maintain portfolio visibility and identify where capacity or specialist involvement is constraining flow. Governance can then become more embedded through environment strategy, policies, platform controls, inventory, and monitoring. Enablement can become more role-specific through champions, expert clinics, learning journeys, reusable test packs, and fusion-team support.

At greater maturity, the system should learn from its own evidence. Risk tiers can be recalibrated when incidents or validation results reveal weak assumptions. Patterns can be promoted, revised, or retired according to observed use. Portfolio measures can change investment priorities. Community feedback can reveal where the safe route is difficult to follow. The mature state is not a large catalog of controls and assets. It is an operating model that can detect when its own design no longer supports the balance it intends to create. These three descriptions are a deliberately simplified view, sufficient for the argument of this paper. The maturity model I set out in Copilot Studio and the Future of the Power Platform Center of Excellence treats the same progression in more detail and across more dimensions, and an organization positioning itself should use that rather than the three states sketched here.

OPERATING TEST

Can a team with a worthwhile idea discover the route, obtain the right support, understand the boundaries, and receive a timely decision?

9. Measuring Whether Innovation Is Real

Measurement is where organizations reveal what they truly value. If the dashboard celebrates the number of agents created, the system will produce agents. If it celebrates speed to prototype, teams will optimize the demonstration. A sustainable innovation system measures movement, outcomes, exposure, and organizational learning together.

Four dimensions of an evidence system




Figure 8. Four dimensions of an evidence system.

Flow: are worthwhile ideas moving?

Flow measures the health of the conversion system. Useful indicators include time from submission to initial orientation, time to controlled validation, time blocked by unresolved decisions, and the proportions of initiatives that progress, pause, or stop at each stage. A healthy system does not maximize progression. It makes decisions with sufficient speed and evidence. A high stop rate during discovery may indicate effective prioritization; a high stop rate after expensive scaling work suggests learning arrived too late.

Value: are outcomes being demonstrated?

Value begins with a baseline and an explicit hypothesis. Measures will vary by use case: cycle time, quality, revenue, cost, capacity released, service experience, risk reduction or strategic option value. The Innovation Hub should not claim every benefit itself, but it should require credible ownership of the measure and maintain portfolio visibility. Where direct financial attribution is difficult, the organization should distinguish observed outcomes, validated estimates and untested assumptions.

Risk: is exposure understood and controlled?

Risk measurement should go beyond incident counts. The absence of reported incidents may indicate good control or poor visibility. Useful measures include critical findings, time to remediation, control exceptions, agents without current owners, overdue reviews, stale knowledge sources, excessive permissions, failed tests, retirement backlog and the proportion of high-risk agents with active monitoring and response plans. Several of those measures are innovation debt made visible. The agents without current owners, the overdue reviews, the stale sources and the retirement backlog are precisely the unfinished decisions

described in chapter 1, and reading them together rather than separately gives the organization a view of how much unresolved work it is carrying. Debt that can be seen can be valued, prioritized, and paid down. Debt that appears only as scattered exceptions accumulates without anyone deciding to accumulate it.

Capability: is the organization learning to repeat success?

Capability measures the compounding effect of innovation. Indicators include reuse of approved components and patterns, breadth of specialist participation, maker confidence, community response time, progression of role-based skills, reduction in recurring defects and the number of improvements to standards or assets generated by delivery learning. Capability is the reason the tenth initiative should be easier and safer than the first.

Most of those indicators describe activity or sentiment, which is a limitation worth naming in a paper that argues adoption means changed work rather than completed training. At least one measure should describe behavior: whether a task is genuinely performed differently now, whether use is sustained after the accompanying support ends, or whether people challenge and correct an agent when the situation calls for it. Those are harder to collect than attendance and they are the only ones that evidence the claim.

Capability measures should also be read by audience rather than in aggregate. Confidence means something different for a maker who has to design an agent and for an end user who has to decide whether to trust one, and a single figure that blends them will move for reasons nobody can interpret. The same applies to reuse, which is a measure of pattern quality among makers and a measure of nothing at all among end users.

Dimension	Leading signals	Lagging signals	Decision enabled
Flow	Intake quality, response time, blocker age.	Stage conversion and end-to-end time.	Where should friction or capacity be addressed?
Value	Baseline quality, user validation, benefit hypothesis.	Realized benefit, sustained use and service outcome.	Which initiatives deserve continued investment or scale?
Risk	Risk-tier accuracy, test coverage, open findings.	Incidents, control failures, overdue reviews and exposure trends.	Where should controls, restrictions, or retirement change?
Capability	Coaching reach, pattern use, community participation.	Reuse, defect reduction, confidence, and distributed delivery capacity.	Which enablement investments improve repeatability?

Table 8. A balanced innovation scorecard

Metrics should drive decisions

A metric without an owner or decision rule becomes reporting theater. The monthly portfolio review should use evidence to reallocate support, remove blockers, challenge value assumptions, increase, or reduce governance intensity, and retire initiatives that no longer justify attention. The quarterly review should examine whether the operating model itself is learning: which policies create repeated friction, which patterns are missing, which skills are scarce, and which strategic themes are producing credible value.

The scorecard above is a starting point for that work rather than a set of measures to adopt. Defining metrics, assigning owners and agreeing decision rules is an exercise each organization has to perform for itself, because the rule is where the judgment lives and it depends on context that no document can supply. What the dimensions offer is the shape of the conversation and a check on its balance: a scorecard that has grown strong on flow and weak on capability is telling the organization something about what it values.

One worked example shows what the exercise produces. The measure is the age of the oldest unresolved blocker in the portfolio. The owner is the person chairing the monthly portfolio review. The rule is that any blocker older than one review cycle is escalated by name to the strategic forum, with the decision it is waiting for stated explicitly. That is a complete measure: something counted, somebody accountable, and a defined consequence. Without the last two parts it is a number on a slide.

A decision-oriented review

A useful portfolio review is organized around decisions rather than presentations. The evidence should be available before the meeting, and exceptions should receive more attention than routine progress. Leaders should be able to see which initiatives need a decision, which assumptions remain unresolved, where exposure has changed, which benefits are supported by evidence, and which work no longer justifies investment. This keeps governance connected to action and prevents reporting from becoming a substitute for management.

The review should also separate initiative performance from system performance. An initiative may be delayed because its hypothesis is weak, because a shared dependency is unavailable, or because the operating model cannot provide a timely decision. Those conditions require different responses. The first may justify stopping. The second may justify investment in a reusable capability. The third may require redesign of governance or enablement. A balanced scorecard becomes useful when it helps leaders distinguish among them.

Finally, decisions should be recorded in a form that can be reused. The record should capture the evidence considered, the accountable role, the decision, any conditions, and the next review trigger. This creates traceability without requiring a long narrative and gives future teams access to the reasoning behind patterns, boundaries, and exceptions.

MEASUREMENT PRINCIPLE

The purpose of measurement is not to prove that innovation activity exists. It is to determine whether that activity is producing value worth sustaining.

Innovation With Direction and Momentum

The agentic era does not make the fundamentals of innovation obsolete. It exposes them. An organization still needs a meaningful problem, a credible hypothesis, capable people, responsible execution, evidence, and an owner prepared to sustain the outcome. What changes is the speed at which possibilities appear, the ease with which prototypes can be created and the potential reach of systems that can interpret context and act.

That combination makes imbalance expensive. Governance without enablement may reduce visible activity while pushing experimentation outside the operating model. Enablement without governance may accelerate creation while multiplying duplication, uncertainty, and exposure. Neither condition produces a system that can repeatedly turn ideas into trusted value.

The answer is not a perfect midpoint. Sustainable innovation requires a dynamic balance. Governance provides strategic direction, proportionate boundaries and accountability. Enablement creates the capability, participation, and momentum to work within those boundaries. Evidence connects them and changes the decision: continue, adapt, scale, or stop.


For a Power Platform and Copilot Studio Center of Excellence, this means operating the Innovation Hub as an organizational integrator rather than as a gatekeeper or a passive support function. The Innovation Hub connects strategy, governance, the Community Hub, delivery, operations, and measurement. It makes the safe route visible and usable. It also ensures that the experience of real initiatives changes future standards, assets, skills, and investment priorities.

The strongest innovation systems will not be those that approve every idea, build the most agents or eliminate all uncertainty. They will be those that learn faster without losing control of what matters. They will know which experiments are worth running, what evidence is required, when a result is ready to scale and when an initiative should end. They will treat stopping as learning, governance as design and enablement as infrastructure for organizational capability.

In the agentic era, the organizations that innovate successfully will not be those that generate the most ideas or impose the most controls. They will be those capable of giving innovation both direction and momentum, and of recognizing through evidence what truly deserves to scale.

Continue the Series

This white paper develops one proposition within the broader body of work I am creating on how organizations design, govern and sustain Centers of Excellence for Microsoft Power Platform and Copilot Studio.



Defining the Framework Structure for the Power Platform Center of Excellence

[View on Amazon ->](#)

<https://www.amazon.com/dp/B0GDDRCD2C>




Copilot Studio and the Future of the Power Platform Center of Excellence

[View on Amazon ->](#)

<https://www.amazon.com/dp/B0H2SXQ3QQ>

Access to these and more material here:

<https://nfernandezba.github.io/Building-the-Center-of-Excellence-for-Power-Platform/>



Sources and Further Reading

The paper uses natural attribution rather than numbered citations. The following sources provide product context, adoption guidance and wider risk-management perspectives. Accessed August 2026.

Microsoft research and organizational guidance

Microsoft WorkLab: 2026 Work Trend Index: Agents, human agency, and the opportunity for every organization
<https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization>

Microsoft Learn: Build an agentic Center of Excellence
<https://learn.microsoft.com/en-us/agents/center-of-excellence/>

Microsoft Learn: Agentic artificial intelligence adoption maturity model: Organization and culture
<https://learn.microsoft.com/en-us/agents/adoption-maturity-model/maturity-model-readiness>

Power Platform and Copilot Studio guidance

Microsoft Learn: Security and governance in Microsoft Copilot Studio
<https://learn.microsoft.com/en-us/microsoft-copilot-studio/security-and-governance>

Microsoft Learn: Build thriving Copilot Studio communities
<https://learn.microsoft.com/en-us/microsoft-copilot-studio/guidance/community-framework-overview>

Microsoft Learn: Define roles and responsibilities for Power Platform adoption
<https://learn.microsoft.com/en-us/power-platform/guidance/adoption/roles>

Microsoft Learn: Power Platform adoption maturity model
<https://learn.microsoft.com/en-us/power-platform/guidance/adoption/maturity-model-details>

Microsoft Learn: What's new in Microsoft Power Platform guidance
<https://learn.microsoft.com/en-us/power-platform/guidance/adoption/whats-new>

Risk and responsible AI

National Institute of Standards and Technology: Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile
<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>

Organization for Economic Co-operation and Development: Artificial Intelligence Principles
<https://www.oecd.org/en/topics/ai-principles.html>

My published framework

Defining the Framework Structure for the Power Platform Center of Excellence

Copilot Studio and the Future of the Power Platform Center of Excellence

Designing the Community Hub

SOURCE BOUNDARY

Product capabilities and platform guidance change over time. Validate implementation decisions against current official documentation and the organization's own legal, security, compliance, and responsible AI requirements.